

INFO 376

Contents

1. Our problem
2. Exploratory analysis
3. Data pipeline (preprocessing, item-item matrix, etc.)
4. Cold-start solution
5. Demo
6. Evaluations and outcomes

Meet the team

Jah Chen

- Cleaning code for readability
- Building embeddings from features
- Preprocessing data into model-friendly matrices

George Lee

- Cleaning Data
- Creation of Recommendation list based on KNN

Aidan Bartlett

- Initial Data ETL
 - Web App
- Cold-start user profiling

Asad Jaffery

- Determine features used for rec system
- Clean and merge data to load for system

Allen Zheng

- Conducted numerous user tests for model evaluation
- Evaluated model efficiency and performance

Amazon Recommender


**Our goal: Build a recommendation web app
to learn about recommender development**

We are using a dataset of Amazon products and reviews,
curated by the UCSD [McAuley Lab](#), to develop our
recommendation engine

We wanted to experiment with developing a
recommendation engine based on real-world data using
what we learned during INFO 376




Project roadmap



Initial Dataset

The more data, the better for us

We chose to use UCSD's data repository, as it gave access to a plethora of data leading to almost 5GB of amazon products and reviews to build our recommendation system from.



Initial Data Cleaning

Filling:

Key numerical values such as price were filled based on the factors such as the median price for a category

Dropping:

High missing cols with little data such as bought_togehter were dropped all together

Textual Representation Creation:

Key text columns such as title, description, user reviews with write ups were joined into a central column for TF-IDF vectorization.

Reviews Data

Missing Values Summary:

	missing_count	missing_%
title	804	0.02
text	1033	0.02
Unnamed: 0	0	0.00
rating	0	0.00
images	0	0.00
asin	0	0.00
parent_asin	0	0.00
user_id	0	0.00
timestamp	0	0.00
helpful_vote	0	0.00
verified_purchase	0	0.00

Review data was only missing a few titles and text content

Products Data

Missing Values Summary:

	missing_count	missing_%
bought_together	137269	100.00
author	137007	99.81
subtitle	136919	99.75
price	75261	54.83
main_category	11036	8.04
store	4375	3.19
title	9	0.01
parent_asin	0	0.00
details	0	0.00
categories	0	0.00
Unnamed: 0	0	0.00
videos	0	0.00
description	0	0.00
features	0	0.00
rating_number	0	0.00
average_rating	0	0.00
images	0	0.00

Over 50% of items lack clear pricing data, and other important components

Pipeline Overview

Aggregate data for each item

Turn text into tf-idf matrix

Turn numerics into sparse csr matrix

Combine both to get the hybrid matrix

Create index for tf-idf

Ready for Recommendations

```
# TF-IDF vectorization for textual metadata
tfidf = TfidfVectorizer(max_features=max_features, min_df=min_df, ngram_range=ngram_range, stop_words='english')
tfidf_matrix = tfidf.fit_transform(item_df['text'])

# Select and normalize numeric features
numeric_features = ['price', 'avg_rating', 'num_ratings']
scaler = StandardScaler()

# Convert to numeric safely
numeric_data = df[numeric_features].apply(pd.to_numeric, errors="coerce").fillna(0)
numeric_scaled = scaler.fit_transform(numeric_data)
print(f"Numeric features scaled (columns: {numeric_features})")

# Convert to sparse format for concatenation
numeric_sparse = np.nan_to_num(numeric_scaled)

# Concatenate text and numeric representations
hybrid_matrix = hstack([tfidf_matrix, numeric_sparse]).tocsr()
print(f"Final hybrid matrix shape: {hybrid_matrix.shape}")

# Maintain item lookup for interpretation
tfidf_index = df['parent_asin'].reset_index(drop=True)
print("Created lookup table linking vectors to item parent_asin.\n")
print(tfidf_index.head())

hybrid_matrix, tfidf, tfidf_index
```

Getting the k-most similar

Problem: Once we get the user's explicit ratings, now how do we actually represent them in the item vector space to get their k-nearest items?

Solution: Create an embedding for each of the items they rated (+1 and +5), in the same way the rest of the catalog was encoded. Then, find the items from the dataset that are most similar to the products that the user rated highly.

```
[('B000068WSA', np.float64(4.0)),
 ('B000HNSKNK', np.float64(4.0)),
 ('B00N5ZDN8Q', np.float64(4.0)),
 ('B001W30G44', np.float64(4.0)),
 ('B0015RCVRM', np.float64(4.0)),
 ('B018K31N68', np.float64(4.0)),
 ('B0001DB6J0', np.float64(4.0)),
 ('B000IN4V6S', np.float64(4.0)),
 ('B00004SV4Y', np.float64(4.0)),
 ('B0030VNLS4', np.float64(4.0))]
```

$$\hat{r}_j = \frac{\sum_{i \in R} r_i \cdot \text{sim}(i, j)}{\sum_{i \in R} |\text{sim}(i, j)|}$$

```
# For each review the user has given:
for review in user_reviews:
    parent_asin = review['parent_asin']
    rating = review['rating']

    # Double check validity
    if parent_asin in tfidf_index.values:

        # Get the row index of the item in the hybrid matrix, then compute cosine similarity
        row_idx = tfidf_index[tfidf_index == parent_asin].index[0]
        sims = cosine_similarity(hybrid_matrix[row_idx], hybrid_matrix)[0]

        # Top_k determines how many similar items to consider for each item the user has rated
        # runtime gets longer for higher k vals, we can discuss val later
        k = top_k

        # Get the indices of the top k similar items
        top_k_idx = np.argpartition(sims, -k)[-k:]
        top_k_idx = top_k_idx[np.argsort(sims[top_k_idx])[:-1]]
        top_k_sims = sims[top_k_idx]

        # calculate those similarity scores
        # formula I use is user rating of current item * similarity score
        for neighbor_idx, sim_val in zip(top_k_idx, top_k_sims):

            # Exclude self-similarity
            if neighbor_idx == row_idx:
                continue
            weight = rating * sim_val
            scores[neighbor_idx] = scores.get(neighbor_idx, 0) + weight
            sim_sums[neighbor_idx] = sim_sums.get(neighbor_idx, 0) + abs(sim_val)

# sort our scores, and also normalize
ranked_scores = {idx: score / sim_sums[idx] for idx, score in scores.items()}
ranked_items = sorted(ranked_scores.items(), key=lambda x: x[1], reverse=True)
```

Challenges we faced

Problem 1: Choosing recommendation system

We decided to use a kNN, item-based content-filtering recommendation engine because our data is naturally **very** sparse, too sparse for matrix factorization or SVD.

Problem 2: Cold start

Providing recommendations to the user given no user history is a classic recommendation problem

Problem 3: Ghost Town

We have no REAL users that has started to use our platform!

Initial User Profile

How should we get likes and dislikes from the user?

Step 1

Get a random, but thoughtful, selection of real products from the dataset

Step 2

Present the products to the user and ask them to rate each one positively or negatively

Step 3

Get the most-similar items to the ones the user rated highly and return them

Sample Dataset


Problem: There are 137,269 unique products; how can we give the user a purposeful but unbiased distribution?

Strategy:



Analyze the dataset and pull out most common groups, then randomly select a sample from those groups.

Common key groups of products:

- Main categories
- Price buckets
- Random
- Overall popularity



Key Product Groups



We also collected the top 10 most-rated products (overall popularity) and a completely random sample



Demo

How our model performed?

	Precision	NDCG
Looking for Racing Games	0.1	0.441
User looking for Mouse and Headset	0.67	0.72
PC gamer looking for game codes	0.60	0.85



Year/Quarter/Month

EVALUATIONS

Our Evaluation Method:

- After users have selected their initial likes and dislikes, recommendations are given to the user
- Users can then provide feedback to whether they like the recommendations that are provided
- We collected data from numerous trial runs focusing on different products and personas each time rating if we liked them or not


Evaluation Results



Evaluation Results

SUMMARY REPORT

Metric	Value
Total Evaluations	59
Total Sessions	11
Total Likes (5.0)	42
Total Dislikes (1.0)	17
Overall Satisfaction Rate (%)	71.19
Overall Dislike Rate (%)	28.81
Average Ratings per Session	5.36
Average Session Satisfaction Rate (%)	67.32
Best Session Satisfaction Rate (%)	100.00
Worst Session Satisfaction Rate (%)	0.00
Unique Products Recommended	54
Products Recommended Multiple Times	5
95% Confidence Interval Lower	59.63
95% Confidence Interval Upper	82.74
Margin of Error (%)	11.56





Any questions?



